

9

Polymorphism and generic programming

What happens when the edict of intention runs up against the edict of irredundancy? The edict of intention calls for expressing clearly the intended types over which functions operate, so that the language can provide help by checking that the types are used consistently. We've heeded that edict, for example, in our definition of the higher-order function `map` from the previous chapter, repeated here:

```
# let rec map (f : int -> int) (xs : int list) : int list =
#   match xs with
#   | [] -> []
#   | hd :: tl -> f hd :: (map f tl) ;;
val map : (int -> int) -> int list -> int list = <fun>
```

The `map` function is tremendously useful for a wide variety of operations over integer lists. It seems natural to apply the same idea to other kinds of lists as well. For instance, we may want to define a function to double all of the elements of a `float` list or implement the `prods` function from Section 7.3.1 to take the products of pairs of integers in a list of such pairs. Using `map` we can try

```
# let double = map (fun x -> 2. *. x) ;;
Error: This expression has type int but an expression was expected
      of type
          float
# let prods = map (fun (x, y) -> x * y) ;;
Error: This pattern matches values of type 'a * 'b
      but a pattern was expected which matches values of type int
```

but we run afoul of the typing constraints on `map`, which can only apply functions of type `int -> int`, and not `float -> float` or `int * int -> int`.

Of course, we can implement a version of `map` for lists of these types as well:

```

# let rec map_float_float (f : float -> float)
#           (xs : float list)
#           : float list =
#   match xs with
#   | [] -> []
#   | hd :: tl -> f hd :: (map_float_float f tl) ;;
val map_float_float : (float -> float) -> float list -> float list
= <fun>

# let rec map_intpair_int (f : int * int -> int)
#           (xs : (int * int) list)
#           : int list =
#   match xs with
#   | [] -> []
#   | hd :: tl -> f hd :: (map_intpair_int f tl) ;;
val map_intpair_int : (int * int -> int) -> (int * int) list -> int
list =
<fun>

```

This is where we run up against the edict of irredundancy: we've written the same code three times now, once for each set of argument types.

What we'd like is a way to map functions over lists *generically*, while still obeying the constraint that whatever type the list elements are, they are appropriate to apply the function to; and whatever type the function returns, the map returns a list of elements of that type.

9.1 Type inference and type variables

The solution to this quandary is found in type inference. In a language with type inference, like OCaml, the type inference process combines all of the type constraints implicit in the use of typed functions together with all of the constraints in explicit typings to compute the types for all of the expressions in a program. For instance, in the definition

```

# let succ x = x + 1 ;;
val succ : int -> int = <fun>

```

it follows from the fact that the `+` function is applied to `x` that `x` must have the same type as the argument type for `+`, that is, `int`. Similarly, since `succ x` is calculated as the output of the `+` function, it must have the same type as `+`'s output type, again `int`. Since `succ`'s argument is of type `int` and output is of type `int`, its type must be `int -> int`. And in fact that is the type OCaml reports for it, even though no

explicit typings were provided.

Propagating type information in this way results in a fully instantiated type `int -> int`. But what if there aren't enough constraints in the code to yield a fully instantiated type? The identity function `id`, which just returns its argument unchanged, is an example:

```
# let id x = x ;;
val id : 'a -> 'a = <fun>
```

Since `x` is never involved in any applications in the definition of `id`, there are no type constraints on it. All that we can conclude is that whatever type `x` is – call it α – the `id` function must take values of type α as argument and return values of type α as output. That is, `id` must be of type $\alpha \rightarrow \alpha$.

The `id` function doesn't have a fully instantiated type. It is a **POLYMORPHIC FUNCTION**, with a **POLYMORPHIC TYPE**. The term *polymorphic* means *many forms*; the `id` function can take arguments of many forms and operate on them similarly.

To express polymorphic types, we need to extend the type expression language. We use **TYPE VARIABLES** to specify that *any* type can be used. We write type variables as identifiers with a prefixed quote mark – `'a`, `'b`, `'c`, and so forth – and conventionally read them as their corresponding Greek letter – α (alpha), β (beta), γ (gamma) – as we've done above. Notice that OCaml has reported a polymorphic type for `id`, namely, `'a -> 'a` (read, " α to α "). This type makes the claim, "for *any* type α , if `id` is applied to an argument of type α it returns a value of type α ."

9.2 Polymorphic map

Returning to the `map` function, we wanted a way to map functions over lists generically. If we just remove the typings in the definition of `map`, it would seem that we could have just such a function, a polymorphic version of `map`.

```
# let rec map f xs =
#   match xs with
#   | [] -> []
#   | hd :: tl -> f hd :: (map f tl) ;;
val map : ('a -> 'b) -> 'a list -> 'b list = <fun>
```

This function performs the same computation as the previous version of `map`, just without any of the explicit type constraints enforced. The function `f` is applied to elements of `xs` and returns elements that appear in the result list, so the type of the argument of `f` must be the type of the elements of `xs` and the type of the result of `f` must be the

type of the elements of the returned list *simply as a consequence of the structure of the code.*

Happily, the type inference process that OCaml uses – developed by Roger Hindley (Figure 9.1) and Robin Milner (Figure 1.7) – infers these constraints automatically, concluding that `map`, like `id`, has a polymorphic type, which the OCaml type inference system has inferred and reported as `('a -> 'b) -> 'a list -> 'b list`. This type expresses the constraint that “**for any types α and β , if `map` is applied to a function from α values to β values, it will return a function that when given a list of α values returns a list of β values.**”

This polymorphic version of `map` can be used to implement `double` and `prods` as above. In each case, the types for these functions are themselves properly inferred by instantiating the type variables of the polymorphic `map` type.

```
# let double = map (fun x -> 2. *. x) ;;
val double : float list -> float list = <fun>
# let prods = map (fun (x, y) -> x * y) ;;
val prods : (int * int) list -> int list = <fun>
```

As inferred by OCaml, `double` takes a `float list` argument and returns a `float list`, and `prods` takes an `(int * int) list` argument and returns an `int list`.

9.3 Regaining explicit types

By taking advantage of polymorphism in OCaml, we've satisfied the edict of irredundancy by defining a polymorphic version of `map`. Unfortunately, we seem to have forgone the edict of intention, since we are no longer explicitly providing information about the intended type for `map`.

But by using the additional expressivity provided by type variables, we can express the intended typing for `map` explicitly.

```
# let rec map (f : 'a -> 'b) (xs : 'a list) : 'b list =
#   match xs with
#   | [] -> []
#   | hd :: tl -> f hd :: (map f tl) ;;
val map : ('a -> 'b) -> 'a list -> 'b list = <fun>
```

The type variables make clear the intended constraints among `f`, `xs`, and the return value `f xs`.

Problem 54  *Source: CS51 2016 midterm; CS51 2017 first midterm*
For each of the following types construct an expression (with no explicit typing annotations, that is, no uses of the `:` operator) for which



Figure 9.1: J. Roger Hindley (1939–), codeveloper with Robin Milner (Figure 1.7) of the Hindley-Milner type inference algorithm that OCaml relies on for inferring the most general polymorphic types for expressions.

*OCaml would infer that type. (The expressions need not be practical or do anything useful; they need only have the requested type.) For example, for the type `bool * bool`, the expression `true, true` would be a possible answer.*

1. `bool * bool -> bool`
2. `'a list -> bool list`
3. `('a * 'b -> 'a) -> 'a -> 'b -> 'a`
4. `int * 'a * 'b -> 'a list -> 'b list`
5. `bool -> unit`
6. `'a -> ('a -> 'b) -> 'b`

□

Exercise 55 Define polymorphic versions of `fold` and `filter`, providing explicit polymorphic typing information. □

Problem 56 Source: CS51 2016 midterm For each of the following definitions of a function `f`, give its most general type (as would be inferred by OCaml) or explain briefly why no type exists for the function.

1. `let f x =
x +. 42. ;;`
2. `let f g x =
g (x + 1) ;;`
3. `let f x =
match x with
| [] -> x
| h :: t -> h ;;`
4. `let rec f x a =
match x with
| [] -> a
| h :: t -> h (f t a) ;;`
5. `let f x y =
match x with
| (w, z) -> if w then y z else w ;;`

□

9.4 The List library

One way, perhaps the best, for satisfying the edict of irredundancy is to avoid writing the same code twice by *not writing the code even once*, instead taking advantage of code that someone else has already written. OCaml, like many modern languages, comes with a large set of libraries (packaged as modules, which we'll cover in (Chapter 12) that provide a wide range of functions. The `List` module in particular provides exactly the higher-order list processing functions presented in this and the previous chapter as polymorphic functions. The documentation for the `List` module gives typings and descriptions for lots of useful list processing functions. For instance, the module provides the map, fold, and filter abstractions of Chapter 8, described in the documentation as

- `map : ('a -> 'b) -> 'a list -> 'b list`
`map f [a1; ...; an]` applies function `f` to `a1, ..., an`, and builds the list `[f a1; ...; f an]` with the results returned by `f`. Not tail-recursive.¹
- `fold_left : ('a -> 'b -> 'a) -> 'a -> 'b list -> 'a`
`fold_left f a [b1; ...; bn]` is `f (... (f (f a b1) b2) ...) bn`.
- `fold_right : ('a -> 'b -> 'b) -> 'a list -> 'b -> 'b`
`fold_right f [a1; ...; an] b` is `f a1 (f a2 (... (f an b) ...))`. Not tail-recursive.
- `filter : ('a -> bool) -> 'a list -> 'a list`
`filter p l` returns all the elements of the list `l` that satisfy the predicate `p`. The order of the elements in the input list is preserved.

They can be invoked as `List.map`, `List.fold_left`, and so forth.

The library provides many other useful functions, including

- `append : 'a list -> 'a list -> 'a list`
 Concatenate two lists. Same as the infix operator `@`. Not tail-recursive (length of the first argument).
- `partition : ('a -> bool) -> 'a list -> 'a list * 'a list`
`partition p l` returns a pair of lists (`l1, l2`), where `l1` is the list of all the elements of `l` that satisfy the predicate `p`, and `l2` is the list of all the elements of `l` that do not satisfy `p`. The order of the elements in the input list is preserved.

¹ We'll come back to the issue of tail recursion in Section 16.2.

The `List` library has functions for sorting, combining, and transforming lists in all kinds of ways.

Although these functions are built into OCaml through the `List` library, it's still useful to have seen how they are implemented and why they have the types they have. In particular, it makes clear that the power of list processing via higher-order functional programming doesn't require special language constructs; they arise from the interactions of simple language primitives like first-class functions and structured data types.

Problem 57 *Source: CS51 2016 midterm* Provide an implementation of the `List.map` function over a list using only a call to `List.fold_right` over the same list, or give a concise argument for why it's not possible to do so. □

Problem 58 *Source: CS51 2016 midterm* Provide an implementation of the `List.fold_right` function using only a call to `List.map` over the same list, or give a concise argument for why it's not possible to do so. □

9.5 Function composition

Source: CS51 2016 midterm The **COMPOSITION** of two unary functions `f` and `g` is the function that applies `f` to the result of applying `g` to its argument.

For example, suppose you're given a list of pairs of integers, where we think of each pair as containing a number and a corresponding weight. We'd like to compute the **WEIGHTED SUM** of the numbers, that is, the sum of the numbers where each has been weighted according to (that is, multiplied by) its weight. Recall the `sum` function from Exercise 39 and the `prods` function from Section 7.3.1. The weighted average of a pair-list can be computed by applying the `sum` function to the result of applying the `prods` function to the list. Thus, `weighted_sum` is just the composition of `sum` and `prods`.

Problem 59 Provide an OCaml definition for a higher-order function `@+` that takes two functions as arguments and returns their composition. The function should have the following behavior:

```
# let weighted_sum = sum @+ prods ;;
val weighted_sum : (int * int) list -> int = <fun>
# weighted_sum [(1, 3); (2, 4); (3, 5)] ;;
- : int = 26
```

Notice that by naming the function `@+`, it is used as an infix, right-associative operator. See [the operator table in the OCaml documentation](#).

tion for further information about the syntactic properties of operators. When defining the function itself, though, you'll want to use it as a prefix operator by wrapping it in parentheses, as `(@+)`. □

Problem 60  What is the type of the `@+` function? □

9.6 Weak type variables

The `List` module provides polymorphic `hd` and `tl` functions for extracting the head and tail of a list.

Exercise 61 What are the types of the `hd` and `tl` functions? See if you can determine them without looking them up. □

These can be composed to allow, for instance, extracting the head of the tail of a list, that is, the list's second item.

```
# let second = List.hd @+ List.tl ;;
val second : '_weak1 list -> '_weak1 = <fun>
```

This definition works,

```
# second [1; 2; 3] ;;
- : int = 2
```

but why did the typing of `second` have those oddly named type variables?

Type variables like `'_weak1` (with the initial underscore) are **WEAK TYPE VARIABLES**, not true type variables. They maintain their polymorphism only temporarily, until the first time they are applied.

Weak type variables arise because in certain situations OCaml's type inference can't figure out how to express the most general types and must resort to this fallback approach.

When a function with these weak type variables is applied to arguments with a specific type, the polymorphism of the function disappears. Having applied `second` to an `int list`, OCaml further instantiates the type of `second` to *only* apply to `int list` arguments, losing its polymorphism. We can see this in two ways, first by checking its type directly,

```
# second ;;
- : int list -> int = <fun>
```

and `second` by attempting to apply it to a list of another type,

```
# second [1.0; 2.1; 3.2] ;;
Error: This expression has type float but an expression was
expected of type
      int
```

To correct the problem, you can of course add in specific typing information

```
# let second : float list -> float =
#   List.hd @+ List.tl ;;
val second : float list -> float = <fun>
```

but this provides no polymorphism. Alternatively, you can provide a full specification of the call pattern in the definition rather than the partial application that was used above:

```
# let second x = (List.hd @+ List.tl) x ;;
val second : 'a list -> 'a = <fun>
```

which gives OCaml sufficient hints to infer types more generally. Of course, in this case, the composition operator isn't really helping. We might as well have defined `second` more directly as

```
# let second x = List.hd (List.tl x) ;;
val second : 'a list -> 'a = <fun>
```

For the curious, if you want to see what's going on in detail, you can check out the discussion in the section “[A function obtained through partial application is not polymorphic enough](#)” in the OCaml FAQ.